

Code Structure Assessment

Directory Structure and Module Organisation

The Townlet repository follows a clear, well-documented structure. Key top-level directories include:

- `configs/` – Contains canonical YAML configuration files for the simulation (defining features, rewards, affordances, perturbations, etc.) ¹. These serve as entry points for different simulation scenarios and settings.
- `docs/` – Houses product/design documentation and planning artifacts (not executable code) ¹, which complement the codebase with design context.
- `scripts/` – Collection of CLI scripts used to run simulations, training sessions, and various tooling tasks ². These Python scripts act as the main entry points for executing the simulation, training harness, data capture, and other utilities (see **Entry Points** below).
- `src/townlet/` – The main Python package implementing the simulator core and services ³. This is structured into sub-packages by domain:
 - `agents/` – Agent state definitions, personality traits, and relationship handling ⁴
 - `console/` – Console (admin/viewer) interface logic and authentication handlers ⁴
 - `config/` – Configuration schemas, loaders, and feature flag validation logic ⁵
 - `core/` – Core simulation loop orchestration (tick loop, main runtime) ⁶
 - `lifecycle/` – Population lifecycle management (agent spawn/exit conditions, population caps) ⁶
 - `observations/` – Observation builders for different observation **variants** (full/hybrid/compact), masking, embeddings ⁷
 - `policy/` – Policy and control logic (scripted controllers, PettingZoo environment wrappers, PPO integration hooks) ⁸
 - `rewards/` – Reward engine and normalization logic (e.g. needs decay, social rewards, reward clipping) ⁹
 - `scheduler/` – Perturbation scheduling, fairness mechanisms, time dilation controls ¹⁰
 - `snapshots/` – Save/load functionality for simulation state and config identity enforcement ¹¹
 - `stability/` – Stability guardrails, canary monitors, and release promotion mechanisms ¹²
 - `telemetry/` – Telemetry publishing (pub/sub observer API), KPI emitters, event schema definitions ¹³
 - `world/` – The grid-world environment model, including spatial grid, affordances (actions), economy, and queues ¹⁴.

Each module is scaffolded with clear responsibilities and integration points, often marked with TODOs for further implementation ¹⁵.

- `src/townlet_ui/` – A supplementary Python package for the **observer dashboard UI** and related client interfaces. This contains the web-based dashboard logic (e.g. `dashboard.py`) and command palette integration for interacting with a running simulation. For example, the `townlet_ui.dashboard` module is invoked by the observer UI script to run a local dashboard server ¹⁶. This separation cleanly isolates core simulation logic (`townlet`) from UI concerns (`townlet_ui`).
- `tests/` – Pytest test suites targeting the configuration loader, core modules, and other components ¹⁷. The presence of this directory indicates an intent for robust test coverage (detailed under **Testing** below).

- `.github/` – Continuous integration and workflow configurations (CI pipelines, issue/pr templates, etc.) ¹⁸.
- `pyproject.toml` – Project configuration file defining dependencies, build settings, linters, and metadata ¹⁸. This makes the package installable (e.g. via pip) and configures tools like **ruff** (linter), **mypy** (type checking), and **pytest** for the project.

Other notable directories include `data/` (for dataset files such as behaviour cloning datasets) and `artifacts/` (for release bundles or milestone artifacts). For instance, `data/bc_datasets/` contains captured trajectories and manifests for behaviour cloning, along with a `versions.json` index ¹⁹. The `artifacts/` folder holds milestone-specific files like acceptance test configurations and release playbooks (e.g. see `artifacts/m5/...` for a Milestone 5 release bundle). These directories are used to store outputs and deployment artifacts related to model training and operational verification, though they are not part of the runtime code per se.

Key Entry Points and Execution Paths

The **main execution entry points** are provided as Python scripts under the `scripts/` directory ². These scripts are intended to be run directly and cover simulation runs, training jobs, and maintenance utilities. Key entry point scripts include:

- `run_simulation.py` – Runs a headless simulation loop given a YAML config. This script loads the config, instantiates the core `SimulationLoop`, and executes a specified number of ticks in the environment ²⁰. It honours the scenario configuration (including the `config_id` and observation variant specified in the YAML) and is useful for non-interactive simulation runs ²¹.
- `run_training.py` – Launches the PPO training harness for agents. It can operate in different modes (replay, rollout, or mixed) as specified by CLI flags, overriding the `training.source` in the config ²¹. This script coordinates either loading replay data or running live rollouts to train the RL policy, depending on the mode, and is the main entry point for training jobs.
- `run_replay.py` – Replays a recorded scenario or dataset through the environment. It uses previously captured trajectories to run the simulation, primarily for validation and analysis. For example, in CI it can be run with a `--validate` flag to ensure replay bundles are consistent with expectations (checking observation variants, social metrics, etc.) ²².
- `observer_ui.py` – Starts the **observer dashboard** UI server connected to a local simulation instance. It launches a web-based dashboard that visualises the simulation state in real-time ²³. Internally, this script loads a config and `SimulationLoop` similar to `run_simulation`, then calls `townlet_ui.dashboard.run_dashboard(...)` to bring up the interactive dashboard interface ²⁰. This allows developers to monitor agents' behaviour on a map, view telemetry streams, and issue console commands via a browser.
- `capture_scripted.py` – Executes scripted agent behaviours to generate and **capture trajectories** for offline datasets. It runs a simulation (often with scripted policies) and saves the sequence of observations/actions to files (e.g. NumPy .npz or JSON manifests) for later analysis or behaviour cloning training ²⁴. This is part of the data generation pipeline for training datasets, capturing replay data along with metadata.
- *(Additional tools)* – The repository includes several other utility scripts for maintenance and analysis. For example, `validate_affordances.py` performs schema and consistency checks on affordance definitions (ensuring no duplicates and valid checksums in `configs/affordances/*` YAML files) ²⁵. There are also scripts like `curate_trajectories.py` to post-process captured trajectories, `promotion_evaluate.py` to evaluate release promotion criteria, `telemetry_watch.py` and `telemetry_summary.py` for analyzing training logs, and

others as referenced in the ops documentation. Each script focuses on a specific operational task or experiment, serving as entry points to those workflows.

These entry point scripts define the **main execution paths** through the system. Generally, they parse CLI arguments, load the appropriate configuration, construct core objects (like the simulation loop or training harness), and then invoke a run method. For example, both `run_simulation.py` and `observer_ui.py` create a `SimulationLoop` from the loaded config and call `loop.run_for(ticks)` or start the dashboard loop²⁰. The training script sets up a `TrainingHarness` (for PPO) and coordinates rollouts vs. replays based on mode. This design cleanly separates runtime scenarios into different scripts, making it easy for developers to invoke different functionalities (simulation vs. training vs. data collection) as needed.

Dependency Management and Build System

Townlet is a Python 3.11+ project and uses a modern **pyproject.toml** for dependency management and build configuration²⁶ ²⁷. All runtime requirements are declared in the `[project.dependencies]` section of `pyproject.toml`. Key dependencies include:

- **numpy** (for numerical computations)²⁸
- **pydantic** (for config and data validation models)²⁸
- **PettingZoo** (for multi-agent RL environment wrappers)²⁹
- **PyYAML** (for YAML config file parsing)³⁰
- **rich** (for rich text/console output, possibly for the CLI or dashboard)³¹
- **typer** (for CLI building, though CLI scripts currently use argparse)³¹

Development and testing tools are specified as an optional **“dev” extras** group. This includes linters and test frameworks like **ruff** (code linter), **mypy** (type checker), **pytest** (test runner), and stub packages³². Installing the project with `pip install -e .[dev]` will bring in these dev dependencies, as noted in the README quickstart³³.

The project uses **Setuptools** as the build system (specified in `pyproject.toml`) and can be built into a wheel or installed in editable mode for development²⁶. The package name is “townlet” (version 0.1.0)³⁴, and source code is under `src/` per the configuration³⁵. There is no separate `requirements.txt` – the pyproject file centrally manages all dependencies and metadata.

The repository also contains configuration for tooling in `pyproject.toml` (such as Ruff rules, Mypy strict settings, and Pytest settings like test path and Pythonpath configuration³⁶ ³⁷). This indicates a focus on code quality and consistency. For example, MyPy is run in strict mode against the `townlet` package³⁸, and Ruff enforces a specific code style³⁹. These tools are likely run in CI to gate contributions (as suggested by the contribution guide and the presence of a `.github/` workflow configuration).

Testing Coverage and Strategy

Testing is handled with **Pytest**, and the project includes a `tests/` directory with various test modules¹⁷. The README suggests running `pytest` to execute a placeholder smoke test suite⁴⁰, meaning a basic set of tests is already in place to sanity-check the setup. According to the repository layout, the tests focus on things like config loading and module integrity¹⁷ – likely ensuring that each module’s scaffold behaves as expected and that configurations parse correctly.

The **overall testing strategy** is documented in the contributing guidelines. It emphasizes a multi-pronged approach to quality ⁴¹ :

- **Unit tests** for core logic such as configuration validation, reward calculations, and lifecycle gate conditions ⁴² . These ensure that individual components (e.g., the reward functions or exit conditions) produce correct results in isolation.
- **Property-based and fuzz tests** for certain systems – for example, tests that validate queue scheduling fairness, economic invariants, and other emergent properties using randomized inputs ⁴² . This helps catch edge cases in complex logic like the scheduler or world interactions.
- **Golden-file tests** for outputs like YAML schemas and telemetry payloads ⁴³ . Golden tests likely compare current output against a known-good output (for example, ensuring that telemetry JSON or saved snapshot formats haven't regressed).
- **Integration smoke tests** targeting critical subsystems. The guide calls out smoke tests for the telemetry transport and observer UI, especially when those areas are modified ⁴⁴ . There are specific test files mentioned (e.g., `test_telemetry_stream_smoke.py`, `test_observer_ui_dashboard.py`, etc.) to exercise end-to-end behaviour of the telemetry pipeline and UI components ⁴⁵ .
- **Snapshot regression tests** for persistence and randomness-sensitive features ⁴⁶ . For instance, tests like `test_snapshot_manager.py` or RNG-related tests ensure that simulation snapshots and random number sequences remain consistent across runs (important for reproducibility).

From the CI configuration (e.g., `tox.ini`), we see a number of test files covering the **world model** (local view, nightly reset, queue integration, rivalry mechanics) and the **console/telemetry systems** ⁴⁷ ⁴⁸ . This indicates that the test suite (while possibly still growing) already covers many core aspects of the simulation: world state progression, console command handling, and telemetry event flows.

Overall, the testing approach shows an intention for thorough coverage: low-level unit tests up to high-level integration tests. Coverage levels are not explicitly stated, but the presence of numerous test files and a requirement to run all tests (with `pytest`) before contributions ⁴⁹ suggests that maintaining and improving test coverage is a priority. The project's testing philosophy is to catch issues early with strict validation (e.g., config sanity checks) and to guard against regressions with targeted smoke tests and golden comparisons.

Configuration Management

Configuration for Townlet is managed through structured **YAML files** in the `configs/` directory ¹ . Each YAML config file encapsulates a simulation setup – including feature flags, environment parameters, reward settings, agent definitions, etc. – and is identified by a `config_id` . The quickstart and docs emphasize aligning the `config_id` with the design specification before running simulations ⁵⁰ . This implies each config file has a field (likely `config_id`) that ties it to a particular scenario or version, ensuring that the running code and design documents stay in sync.

At runtime, configs are loaded and validated by the `townlet.config.loader` module. This loader uses **PyYAML** to parse the YAML and **Pydantic** models to validate and instantiate configuration objects ⁵¹ . The code centralises config parsing and applies sanity checks. For example, it enforces valid combinations of settings and applies **feature flags** to toggle experimental systems ⁵² . The configuration schema (defined via Pydantic `BaseModel` classes) covers various sections such as feature toggles (`FeatureFlags` for enabling/disabling stages or systems), reward function parameters (`RewardsConfig` with constraints on values), scheduler settings, etc., reflecting the

requirements in the design docs ⁵² ⁵³. This strict schema helps catch misconfigurations early (e.g., the loader will raise validation errors if a value is out of allowed range or a required field is missing).

Environment variable handling is also part of configuration management for sensitive or deployment-specific values. For instance, the console authentication tokens can be specified via environment variables: the config can include a `token_env` name instead of a raw token, and the loader will fetch the actual token from the environment at runtime ⁵⁴. This design allows secrets or environment-specific credentials to be injected without hardcoding them in config files. Another example is the affordance hook allowlist – an environment variable `TOWNLET_AFFORDANCE_HOOK_MODULES` can be used to permit certain dynamic modules, though the production config can disable env-based hooks for security ⁵⁵. Overall, the configuration system is built to be flexible (YAML-driven), while enforcing rules and supporting feature flags to safely turn features on/off as needed.

The repository provides several **example config files** under subfolders like `configs/examples/` and `configs/demo/` (e.g. `poc_hybrid.yaml`, `poc_demo.yaml`) to illustrate various scenario setups. These serve as starting points for running the simulation in different modes (for example, a proof-of-concept demo vs. a hybrid observation mode scenario). Users are expected to pick or modify a YAML config and then run the appropriate script with `--config <file>` to launch that scenario. By keeping configurations in versioned files, Townlet makes it easier to reproduce runs and adjust parameters without changing code.

Build and Deployment Artifacts

As of this audit, Townlet appears to be primarily a source-distributed Python project (version 0.1.0) meant to be installed via pip or run from source. **Build artifacts** in the traditional sense (like compiled binaries or Docker images) are not present in the repository – there is no Dockerfile or container orchestration config included, and the deployment is likely handled by running the Python code directly on a suitable environment. The focus is on packaging the code as a Python library (via the pyproject configuration) and using CLI scripts for execution. This means the “build” is mostly about ensuring the Python package can be installed; `setuptools` is configured accordingly to find the packages in `src/` ³⁵. Developers can produce a wheel or install an editable package, but there isn't a separate build pipeline producing runtime artifacts beyond the Python package itself.

Instead, the repository provides **operational artifact directories** that capture the outputs and resources of running the system, especially for release and evaluation purposes. The `artifacts/` directory contains structured subfolders for different milestones and phases of the project. For example, under `artifacts/m5/release_bundle/` there are playbooks and configuration files related to the Milestone 5 release. These might include acceptance test configurations (`config_idle_v1.yaml` as referenced in docs), promotion drill outputs, and other files needed to reproduce or validate a release. Similarly, `artifacts/phase4/` contains materials for an earlier phase. These artifacts are not build outputs in the traditional sense, but they are important for deployment verification – they bundle the necessary configs, logs, and results to be delivered with a release. In a recovery scenario, these give clues about how the system was tested and what data accompanies a deployment.

The presence of a `data/` directory also points to deployment-related data artifacts. Specifically, `data/bc_datasets/` holds behaviour cloning datasets and metadata that the simulation/training might rely on. This includes captured trajectory files (`captures/`), curated manifests of those trajectories (`manifests/`), and a `versions.json` that tracks dataset versions and their checksums ⁵⁶. These data files would be considered deployment artifacts in the sense that any environment

running Townlet's training pipeline needs these datasets present. The project even provides scripts (`audit_bc_datasets.py`) to verify dataset integrity via checksum, indicating these data artifacts are part of the release quality control ⁵⁷.

In terms of **deployment process**, there is evidence of rigorous operational practices (though not a one-click deployment script). The CI/CD configuration in `.github/` likely runs tests and linters on each push (ensuring the code is clean and passes all checks). The ops documentation (in `docs/ops/` and `docs/rollout/`) outlines manual or semi-automated steps to promote a release, which include running the various scripts mentioned earlier and collecting their outputs into the `artifacts/` bundle ⁵⁸ ⁵⁷. For example, prior to a release, one would run simulation rollouts and training with `run_training.py`, generate telemetry and summary reports (`telemetry_watch.py`, `telemetry_summary.py`), and ensure no regression flags are raised, packaging those results as evidence. This suggests that **deployment** in Townlet's context is tied to milestone-based releases where a collection of configs and result logs are delivered, rather than deploying a live service.

In summary, the project's build and deployment artifacts are mostly **configuration and data-centric**. The codebase itself is packaged via standard Python tooling, and no containerization is provided out-of-the-box. Deployment seems to involve running the Python package on a server or workstation with the required configs and then using the provided scripts to simulate, train, or evaluate as needed. The included artifact directories and data files ensure that anyone recovering the project can find the necessary pieces (configs, datasets, and documentation) to assemble a working environment and verify the system's behaviour in line with its last known state.

Sources:

- Townlet Repository README (project layout and status) ⁵⁹ ¹⁷ ⁴⁰
- `pyproject.toml` (dependencies and tooling configs) ⁶⁰ ³²
- Townlet Contributing Guide (testing expectations and development workflow) ⁴¹ ⁴⁹
- Operations Handbook (entry point usage and deployment workflows) ²¹ ²² ²⁵
- Source code (config loader and console auth for env var usage) ⁵² ⁵⁴
- Dataset README (structure of data and artifact usage) ⁵⁶

¹ ² ³ ⁴ ⁵ ⁶ ⁷ ⁸ ⁹ ¹⁰ ¹¹ ¹² ¹³ ¹⁴ ¹⁵ ¹⁷ ¹⁸ ³³ ⁴⁰ ⁵⁰ ⁵⁹ README.md

<https://github.com/tachyon-beep/townlet/blob/ce19ff3ed21fa8e36cb74c402d72b52eeb93a8c4/README.md>

¹⁶ ²⁰ ²³ `observer_ui.py`

https://github.com/tachyon-beep/townlet/blob/ce19ff3ed21fa8e36cb74c402d72b52eeb93a8c4/scripts/observer_ui.py

¹⁹ ⁵⁶ README.md

https://github.com/tachyon-beep/townlet/blob/ce19ff3ed21fa8e36cb74c402d72b52eeb93a8c4/data/bc_datasets/README.md

²¹ ²² ²⁴ ²⁵ ⁵⁵ ⁵⁷ ⁵⁸ OPS_HANDBOOK.md

https://github.com/tachyon-beep/townlet/blob/ce19ff3ed21fa8e36cb74c402d72b52eeb93a8c4/docs/ops/OPS_HANDBOOK.md

²⁶ ²⁷ ²⁸ ²⁹ ³⁰ ³¹ ³² ³⁴ ³⁵ ³⁶ ³⁷ ³⁸ ³⁹ ⁶⁰ `pyproject.toml`

<https://github.com/tachyon-beep/townlet/blob/ce19ff3ed21fa8e36cb74c402d72b52eeb93a8c4/pyproject.toml>

⁴¹ ⁴² ⁴³ ⁴⁴ ⁴⁵ ⁴⁶ ⁴⁹ CONTRIBUTING.md

<https://github.com/tachyon-beep/townlet/blob/ce19ff3ed21fa8e36cb74c402d72b52eeb93a8c4/CONTRIBUTING.md>

47 48 **tox.ini**

<https://github.com/tachyon-beep/townlet/blob/ce19ff3ed21fa8e36cb74c402d72b52eeb93a8c4/tox.ini>

51 52 53 **loader.py**

<https://github.com/tachyon-beep/townlet/blob/ce19ff3ed21fa8e36cb74c402d72b52eeb93a8c4/src/townlet/config/loader.py>

54 **auth.py**

<https://github.com/tachyon-beep/townlet/blob/ce19ff3ed21fa8e36cb74c402d72b52eeb93a8c4/src/townlet/console/auth.py>